



Gaussian Discriminant Analysis (GDA) - Complete Breakdown

1. In-Depth and Specific Intuitive Understanding

What is Gaussian Discriminant Analysis (GDA)?

Gaussian Discriminant Analysis (GDA) is a **probabilistic classification algorithm** based on **Bayes' theorem** and assumes that each class follows a **Gaussian (Normal) distribution** in feature space.

Instead of directly modeling $P(y|x)$ like Logistic Regression, GDA models the **joint probability** $P(x, y)$ and then applies Bayes' rule to obtain $P(y|x)$.

Key Idea

-

GDA assumes that each class has a **Gaussian distribution**:

$$p(x|y = k) \sim \mathcal{N}(\mu_k, \Sigma)$$

- Using Bayes' Theorem, the probability of class $y = k$ given x is:

$$P(y = k|x) = \frac{P(x|y = k)P(y = k)}{P(x)}$$

- We classify x into the class with the highest posterior probability.

Intuition Behind GDA

- Imagine you have two groups of points (e.g., two species of flowers) and each group follows a **Gaussian (bell-shaped) distribution**.
 - GDA finds the **probability density function** (PDF) of each group.
 - When a new point arrives, GDA **compares probabilities** and assigns the class with the **highest likelihood**.
-

2. When GDA is Used and When It Should Be Avoided

✓ When to Use GDA

- When **features follow a roughly Gaussian distribution**.
- When **data is well-separated** (GDA models decision boundaries well).
- When the number of **samples is small** (parametric models perform better in small-data regimes).
- When **computational efficiency is needed** (GDA has a closed-form solution).

✗ When to Avoid GDA

- If the **Gaussian assumption does not hold**, GDA performs poorly.
 - If the dataset is **high-dimensional**, estimating the covariance matrix Σ is computationally expensive.
 - If the **classes have highly non-linear boundaries**, GDA's linear or quadratic decision boundary may not work well.
-

3. When It Fails to Converge and How to Avoid That

When GDA Fails

- **Singular Covariance Matrix Σ (Non-Invertibility):**
 - Happens when there are **fewer data points than features**.
 - Happens when **some features are perfectly correlated** (multicollinearity).
- **Poor Performance if Data is Not Gaussian:**
 - If the data distribution is skewed or multimodal, GDA struggles.
- **Unstable Estimation in Small Datasets:**
 - Estimating the covariance matrix with very few samples leads to poor generalization.

How to Ensure Convergence

- ✓ **Use Regularized Covariance Estimation (Shrinkage Methods)** for stability.
- ✓ **Reduce feature dimensionality** (e.g., PCA) if $n \ll m$.
- ✓ **Check for Gaussianity** (use normality tests like Shapiro-Wilk).

When GDA Always Converges

- If there are **enough samples** relative to feature dimensions.
 - If the covariance matrix is **well-conditioned** (not singular).
 - If features are **reasonably Gaussian**.
-

4. Advantages and Disadvantages

Advantages

- ✓ Simple and computationally efficient.
- ✓ Works well for small datasets (unlike deep learning models).
- ✓ Interpretable—explicit probability distributions.
- ✓ Handles missing values well (if modeled probabilistically).

Disadvantages

- ✗ Assumes Gaussian distributions, which may not hold in real-world data.
 - ✗ Covariance estimation is expensive for high-dimensional data.
 - ✗ Linear Decision Boundaries (if using shared Σ) may not capture complex patterns.
-

5. Intuitive Algorithm / Pseudo Code

1. Estimate class priors $P(y)$.
2. Estimate mean μ_k and covariance Σ for each class.
3. Compute the likelihood $P(x|y)$ using the Gaussian density function.
4. Apply Bayes' theorem to compute $P(y|x)$.
5. Assign x to the class with the highest posterior probability.

plaintext

 Copy

 Edit

```
1. Compute class priors:  $P(y=k) = (\text{\# samples in class } k) / (\text{total samples})$ 
2. Compute class means:  $\mu_k = \text{mean}(X \mid y=k)$ 
3. Compute shared covariance matrix  $\Sigma$ :
    $\Sigma = (1/m) * \sum (x_i - \mu_{yi}) (x_i - \mu_{yi})^T$ 
4. Compute class-conditional probabilities using Gaussian PDF
5. Use Bayes' theorem to classify new data points
```

6. Mathematical and Logical Breakdown

Step 1: Estimate Class Priors

The probability of each class:

$$P(y = k) = \frac{\text{Number of points in class } k}{\text{Total number of points}}$$

Step 2: Estimate Mean and Covariance

For each class $y = k$, estimate the mean vector:

$$\mu_k = \frac{1}{N_k} \sum_{i:y_i=k} x_i$$

And the shared covariance matrix:

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{y_i})(x_i - \mu_{y_i})^T$$

Step 3: Compute Class-Conditional Probability (Gaussian Density Function)

For a new data point x , the likelihood under class k is:

$$P(x|y = k) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu_k)^T \Sigma^{-1} (x - \mu_k) \right)$$

Step 4: Compute Posterior Probability (Bayes' Rule)

Using Bayes' Theorem:

$$P(y = k|x) = \frac{P(x|y = k)P(y = k)}{P(x)}$$

7. Manual Implementation in Python


```
import numpy as np
```

```
class GaussianDiscriminantAnalysis:
```

```
    def __init__(self):
```

```
        self.classes = None
```

```
        self.means = {}
```

```
        self.covariance = None
```

```
        self.priors = {}
```

```
    def fit(self, X, y):
```

```
        """Train GDA by estimating class priors, means, and covariance matrix."""
```

```
        self.classes = np.unique(y)
```

```
        m, n = X.shape
```

```
        self.covariance = np.zeros((n, n))
```

```
        for c in self.classes:
```

```
            X_c = X[y == c]
```

```
            self.means[c] = np.mean(X_c, axis=0)
```

```
            self.priors[c] = X_c.shape[0] / m
```

```
            self.covariance += np.cov(X_c, rowvar=False) * (X_c.shape[0] - 1)
```

```
        self.covariance /= m
```

```
    def predict(self, X):
```

```
        """Predict class labels using Bayes' theorem."""
```

```
        predictions = []
```

```
        inv_cov = np.linalg.inv(self.covariance)
```

```
        for x in X:
```

```
            posteriors = {}
```

```
            for c in self.classes:
```

```
                mean_diff = x - self.means[c]
```

```
                likelihood = -0.5 * mean_diff.T @ inv_cov @ mean_diff
```

```
                posteriors[c] = np.log(self.priors[c]) + likelihood
```

```
            predictions.append(max(posterior, key=posterior.get))
```

```
        return np.array(predictions)
```

8. Scikit-Learn Implementation (Fully Commented)

python

Copy

Edit

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.model_selection import train_test_split
from sklearn.datasets import make_classification

# Generate synthetic dataset
X, y = make_classification(n_samples=100, n_features=2, n_classes=2, random_state=4)

# Split into training (80%) and testing (20%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)

# Initialize GDA model
model = LinearDiscriminantAnalysis()
model.fit(X_train, y_train)

# Predict class labels
y_pred = model.predict(X_test)

# Print accuracy
print("Model Accuracy:", model.score(X_test, y_test))
```

Final Summary

- GDA is a probabilistic classifier that assumes Gaussian distributions per class.
 - It estimates class priors, means, and a covariance matrix to classify points.
 - Works well when data is Gaussian, but fails in non-Gaussian cases.
 - Computationally efficient for small datasets.
-