



📖 Memory updated

# K-Means Clustering – A Deep Dive

---

# 1. What is Clustering?

Clustering is an **unsupervised learning technique** that groups data points into **clusters** based on similarity.

- Unlike supervised learning, **there are no labels** in the dataset.
  - The goal is to **discover hidden structures** in the data.
-

## 2. What is K-Means?

K-Means is one of the simplest and most widely used clustering algorithms. It works by:

1. Choosing **K cluster centroids** (randomly or heuristically).
  2. Assigning each data point to the **nearest centroid**.
  3. Recomputing centroids as the **average of assigned points**.
  4. Repeating until convergence (centroids stop changing).
- 

## 3. When to Use & Avoid K-Means

### ✅ Use K-Means When:

- You **need fast and scalable clustering**.
- The clusters are **well-separated and convex**.
- You have a **moderate** number of clusters.

### ❌ Avoid K-Means When:

- The dataset has **non-spherical clusters** (e.g., concentric circles).
  - There are **many outliers**, as K-Means is sensitive to them.
  - The **value of K is unknown**, and choosing it is difficult.
-

## 4. Choosing K (Number of Clusters)

### Methods to Determine the Optimal K

#### 1. Elbow Method

- Compute the **within-cluster sum of squares (WCSS)** for different values of K.
- Look for an **"elbow point"** where WCSS stops decreasing significantly.

2.

#### Silhouette Score

- Measures how well-separated the clusters are.
- Higher scores indicate better clustering.

3.

#### Gap Statistic

- Compares clustering performance against randomly generated datasets.
-

# Mathematical Breakdown of Methods for Choosing $K$ in K-Means Clustering

Selecting the optimal number of clusters  $K$  is crucial for interpretable and efficient clustering.

Here, we cover the most common techniques:

---

# 1. The Elbow Method

## Intuition

- We compute the **within-cluster sum of squares (WCSS)** for different values of  $K$ .
- Plot WCSS vs.  $K$  and look for an "elbow" point, where WCSS stops decreasing significantly.

## Mathematics

The **WCSS (Within-Cluster Sum of Squares)** measures how compact each cluster is:

$$\text{WCSS} = \sum_{i=1}^K \sum_{x_j \in C_i} \|x_j - \mu_i\|^2$$

where:

- $x_j$  = data point
- $\mu_i$  = centroid of cluster  $C_i$
- $\|x_j - \mu_i\|^2$  = squared Euclidean distance from  $x_j$  to its centroid

## Procedure

1. Compute WCSS for different values of  $K$ .
2. Plot  $K$  vs. WCSS.
3. Identify the "elbow point" where WCSS reduction slows.

## Limitations

- The elbow point is **subjective**.
- Doesn't always work well if clusters are not clearly defined.

# 1. The Elbow Method

## Intuition

- We compute the **within-cluster sum of squares (WCSS)** for different values of  $K$ .
- Plot WCSS vs.  $K$  and look for an "elbow" point, where WCSS stops decreasing significantly.

## Mathematics

The WCSS (Within-Cluster Sum of Squares) measures how compact each cluster is:

$$\text{WCSS} = \sum_{i=1}^K \sum_{x_j \in C_i} \|x_j - \mu_i\|^2$$

where:

- $x_j$  = data point
- $\mu_i$  = centroid of cluster  $C_i$
- $\|x_j - \mu_i\|^2$  = squared Euclidean distance from  $x_j$  to its centroid

## Procedure

1. Compute WCSS for different values of  $K$ .
2. Plot  $K$  vs. WCSS.
3. Identify the "elbow point" where WCSS reduction slows.

## Limitations

- The elbow point is **subjective**.
- Doesn't always work well if clusters are not clearly defined.



## 2. Silhouette Score

### Intuition

- Measures how well-separated clusters are.
- The score is between **-1 and 1**:
  - 1** → Perfect clustering.
  - 0** → Overlapping clusters.
  - 1** → Wrong clustering.

### Mathematics

For each data point  $x_i$ , compute:

1.

**Average intra-cluster distance** (how close is  $x_i$  to its own cluster centroid?):

$$a(i) = \frac{1}{|C_k| - 1} \sum_{x_j \in C_k, i \neq j} \|x_i - x_j\|$$

where  $C_k$  is the cluster to which  $x_i$  belongs.

2.

**Average inter-cluster distance** (how far is  $x_i$  from the nearest different cluster?):

$$b(i) = \min_{C_m \neq C_k} \frac{1}{|C_m|} \sum_{x_j \in C_m} \|x_i - x_j\|$$

where  $C_m$  is the nearest cluster.

3. Compute **Silhouette Score** for  $x_i$ :

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

4. Compute the **average silhouette score** for all points to evaluate clustering quality.

## Procedure

1. Compute silhouette scores for different  $K$ .
2. Choose  $K$  with the **highest average silhouette score**.

## Limitations

- Requires pairwise distance computations → **Computationally expensive** for large datasets.
- Can be misleading if clusters have non-convex shapes.

### 3. Gap Statistic

#### Intuition

- Compares clustering performance against **randomly generated datasets**.
- A higher gap means clustering is **better than random noise**.

#### Mathematics

1. Compute the WCSS for **real data**:

$$W_k = \sum_{i=1}^K \sum_{x_j \in C_i} \|x_j - \mu_i\|^2$$

2. Generate **B random datasets** (same size as real data, but uniformly distributed).
3. Compute the WCSS for each random dataset:

$$W_k^b = \sum_{i=1}^K \sum_{x_j \in C_i^b} \|x_j - \mu_i^b\|^2, \quad b = 1, \dots, B$$

4. Compute the expected WCSS over all **B random datasets**:

$$E(W_k) = \frac{1}{B} \sum_{b=1}^B W_k^b$$

5. Compute the **Gap Statistic**:

$$G_k = \frac{1}{B} \sum_{b=1}^B \log(W_k^b) - \log(W_k)$$

- 6.

Choose  $K$  where  $G_k$  is maximum.

## Procedure

1. Compute  $G_k$  for different  $K$ .
2. Choose the **largest**  $G_k$ .

## Limitations

- **Computationally expensive** due to multiple random dataset generations.
- Assumes clusters are compact and spherical.

## 4. Davies-Bouldin Index

### Intuition

- Measures **inter-cluster similarity** (how well-separated clusters are).
- Lower values indicate **better clustering**.

### Mathematics

1.

Compute cluster dispersion:

$$S_i = \frac{1}{|C_i|} \sum_{x_j \in C_i} \|x_j - \mu_i\|$$

2. Compute cluster separation between clusters  $i$  and  $j$ :

$$R_{ij} = \frac{S_i + S_j}{\|\mu_i - \mu_j\|}$$

3. Compute **Davies-Bouldin Index**:

$$DB = \frac{1}{K} \sum_{i=1}^K \max_{j \neq i} R_{ij}$$

4.

Choose  $K$  with the **smallest DB index**.

### Procedure

- Compute DB index for different  $K$ .
- Select  $K$  where DB is **minimum**.

### Limitations

- Works well for convex clusters but fails for non-spherical clusters.
- Sensitive to noise.



# Comparison of Methods for Choosing K

Method	Works Best When...	Computational Complexity	Key Limitation
Elbow Method	Clusters are well-separated	Low	The "elbow" is subjective
Silhouette Score	Clusters have clear boundaries	High	Expensive for large datasets
Gap Statistic	Random reference comparison is needed	Very High	Computationally expensive
Davies-Bouldin Index	Clusters have equal density	Medium	Fails for irregular cluster shapes

---

## Final Takeaways

- ✓ Elbow Method is the simplest approach but is subjective.
- ✓ Silhouette Score measures separation quality, but is expensive.
- ✓ Gap Statistic provides a rigorous test but is very slow.
- ✓ Davies-Bouldin Index works well for convex clusters but fails for complex shapes.

## 5. Step-by-Step Pseudo Code for K-Means

plaintext

Copy

Edit

```
# Step 1: Initialize Clusters
1. Choose K (number of clusters).
2. Randomly select K data points as initial centroids.

# Step 2: Assign Points to Clusters
1. For each data point:
    a. Compute distance to each centroid.
    b. Assign point to the nearest centroid.

# Step 3: Recompute Centroids
1. For each cluster:
    a. Compute the mean of all assigned points.
    b. Update the centroid.

# Step 4: Check for Convergence
1. If centroids do not change significantly, stop.
2. Otherwise, repeat Steps 2 and 3.

# Step 5: Output Final Clusters
1. Return cluster assignments and final centroids.
```

## 6. Mathematical Breakdown

### 6.1 Objective Function (Minimizing Variance Within Clusters)

The goal of K-Means is to minimize the **intra-cluster variance**, measured by the sum of squared differences:

$$J = \sum_{i=1}^K \sum_{x_j \in C_i} \|x_j - \mu_i\|^2$$

where:

- $x_j$  is a data point.
- $\mu_i$  is the centroid of cluster  $C_i$ .
- The inner sum computes the variance within a cluster.

### 6.2 Distance Calculation

The most common distance metric used is **Euclidean distance**:

$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

### 6.3 Updating Centroids

The new centroid  $\mu_i$  is simply the **mean** of all assigned points:

$$\mu_i = \frac{1}{|C_i|} \sum_{x_j \in C_i} x_j$$



## 7. Fully Commented Manual Implementation (From Scratch)

```

37     def update_centroids(self, X, labels):
38         """
39         Updates centroids by computing the mean of assigned points.
40         :param X: Data points.
41         :param labels: Cluster assignments.
42         """
43         new_centroids = np.array([X[labels == i].mean(axis=0) for i in range(self.k)])
44         return new_centroids
45
46     def fit(self, X):
47         """
48         Run the K-Means clustering algorithm.
49         :param X: Data points.
50         """
51         self.initialize_centroids(X)
52
53         for i in range(self.max_iters):
54             labels = self.assign_clusters(X) # Step 2: Assign points
55             new_centroids = self.update_centroids(X, labels) # Step 3: Update centroids
56
57             # Check for convergence (if centroids do not change significantly)
58             if np.linalg.norm(self.centroids - new_centroids) < self.tol:
59                 break
60
61             self.centroids = new_centroids # Update centroids
62
63     def predict(self, X):
64         """
65         Assign new data points to the nearest cluster.
66         :param X: New data points.
67         :return: Cluster assignments.
68         """
69         return self.assign_clusters(X)

```

```

1 import numpy as np
2
3 class KMeans:
4     """
5     Implementation of K-Means Clustering Algorithm.
6     """
7     def __init__(self, k=3, max_iters=100, tol=1e-4):
8         """
9         Initialize K-Means parameters.
10        :param k: Number of clusters.
11        :param max_iters: Maximum number of iterations.
12        :param tol: Convergence tolerance (stop if centroid shifts are below this value).
13        """
14        self.k = k
15        self.max_iters = max_iters
16        self.tol = tol # Threshold for stopping criterion
17        self.centroids = None # Stores cluster centroids
18
19    def initialize_centroids(self, X):
20        """
21        Randomly initialize K cluster centroids from the data points.
22        :param X: Data points (numpy array of shape (n_samples, n_features)).
23        """
24        np.random.seed(42) # For reproducibility
25        indices = np.random.choice(X.shape[0], self.k, replace=False)
26        self.centroids = X[indices]
27
28    def assign_clusters(self, X):
29        """
30        Assigns each data point to the nearest centroid.
31        :param X: Data points.
32        :return: Array of cluster assignments.
33        """
34        distances = np.linalg.norm(X[:, np.newaxis] - self.centroids, axis=2) # Compute distances
35        return np.argmin(distances, axis=1) # Assign to the nearest centroid
36

```

## 8. Scikit-Learn Implementation

python

 Copy

 Edit

```
from sklearn.cluster import KMeans
from sklearn.datasets import make_blobs

# Generate synthetic dataset
X, _ = make_blobs(n_samples=300, centers=3, cluster_std=1.0, random_state=42)

# Train K-Means
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans.fit(X)

# Get cluster assignments and centroids
labels = kmeans.labels_
centroids = kmeans.cluster_centers_

print("Cluster Centers:\n", centroids)
```

## 9. Advantages & Disadvantages of K-Means

### ✓ Advantages

- Simple and easy to implement.
- Scalable to large datasets.
- Fast because it uses simple distance calculations.

### ✗ Disadvantages

- Requires manual selection of  $K$ .
  - Fails for non-spherical clusters.
  - Sensitive to outliers.
-

---

## Final Takeaways

- ✓ K-Means is a fundamental clustering algorithm that partitions data into K groups.
- ✓ Choosing the right K is crucial for meaningful clustering.
- ✓ Sensitive to initialization, but optimizations (K-Means++) can improve results.

# Comparison of Popular Clustering Algorithms

Algorithm	Intuition (How It Works & What It Does)	When to Use	Pros	Cons
K-Means	Assigns points to the nearest centroid, then updates centroids iteratively to minimize variance.	When clusters are <b>spherical, well-separated</b> , and you know <b>K</b> .	Simple, fast, works well on large datasets.	Sensitive to <b>initialization</b> and <b>outliers</b> . Struggles with non-spherical clusters.
Hierarchical Clustering	Recursively merges (agglomerative) or splits (divisive) clusters based on similarity.	When you want a <b>dendrogram</b> (cluster hierarchy) and <b>don't know K</b> .	No need to predefine K, creates a full hierarchy.	Slow for large datasets, sensitive to <b>distance metrics</b> .
DBSCAN (Density-Based Clustering)	Groups points that have a minimum number of neighbors within a radius, marking outliers.	When clusters have <b>arbitrary shapes</b> and contain <b>noise</b> .	Detects <b>arbitrary-shaped</b> clusters, robust to noise.	Struggles with varying densities and <b>high-dimensional data</b> .
Gaussian Mixture Model (GMM)	Assumes data is generated from <b>multiple Gaussian distributions</b> and assigns probabilities to clusters.	When clusters <b>overlap</b> and have <b>elliptical shapes</b> .	Soft clustering, flexible cluster shapes.	Computationally expensive, requires <b>Gaussian assumption</b> .
Mean Shift	Moves data points toward areas of <b>highest density</b> (mode seeking).	When clusters have <b>varying densities</b> and you want automatic cluster detection.	No need to specify K, works for <b>irregular densities</b> .	Computationally expensive, sensitive to bandwidth selection.
Spectral Clustering	Uses graph-based methods and eigenvalues to separate data points into clusters.	When clusters are <b>connected but not necessarily spherical</b> . <div></div>	Works well for <b>non-convex clusters</b> , handles complex structures.	Doesn't scale well, computationally expensive.

## Final Takeaways

- ✓ K-Means → Best for simple, well-separated clusters.
  - ✓ DBSCAN → Best for arbitrary-shaped clusters with noise.
  - ✓ Hierarchical Clustering → Best when you don't know K and want a hierarchy.
  - ✓ GMM → Best when clusters overlap and have elliptical shapes.
  - ✓ Mean Shift → Best when you want automatic cluster detection without K.
  - ✓ Spectral Clustering → Best for complex structures and non-spherical clusters.
-