



Statistical Methods in Machine Learning – A Deep Dive

Method	What It Does	When to Use It	Example Use Cases
Maximum Likelihood Estimation (MLE)	Finds parameters that maximize the likelihood of observed data.	When you assume iid data and want a frequentist estimate .	Estimating parameters of a Gaussian distribution , logistic regression.
Maximum A Posteriori (MAP)	Similar to MLE but incorporates a prior belief.	When prior knowledge is available, and you want a Bayesian approach .	Small datasets, Bayesian Logistic Regression , LDA topic modeling .
Softmax Function	Converts raw scores into probabilities over multiple classes.	When dealing with multi-class classification .	Final activation function in Neural Networks , image classification.
Cross-Entropy Loss	Measures the difference between two probability distributions.	When training classification models , especially with Softmax output .	Loss function in Deep Learning , NLP models like BERT .
KL Divergence	Measures how different two probability distributions are.	When comparing two distributions in Bayesian models or RL .	Used in Variational Autoencoders (VAEs) , Bayesian Inference .
Jensen-Shannon Divergence	Symmetric version of KL Divergence.	When comparing two probability distributions in a more stable way .	Used in GANs (Generative Adversarial Networks) .
Bayesian Inference	Updates probability estimates as new data arrives.	When working with uncertain, small, or streaming data .	Used in Naive Bayes classifiers , Bayesian Neural Networks .
Gaussian Mixture Models (GMMs)	Models complex distributions using multiple Gaussians.	When clustering data where each cluster follows a Gaussian distribution .	Used in unsupervised learning , anomaly detection , speech processing .
Expectation-Maximization (EM Algorithm)	Finds the best parameters when there are hidden/missing variables.	When working with latent variable models .	GMMs , HMMs (Hidden Markov Models) , unsupervised learning .



1. Maximum Likelihood Estimation (MLE)

Intuition: What Problem Does MLE Solve?

MLE is a method for estimating parameters of a probability distribution that **maximize the likelihood of observed data**.

Example: Coin Flip

Imagine flipping a biased coin n **times**, and we want to estimate the probability p of landing heads.

Given outcomes $X = [H, T, H, H, T]$, the best estimate of p should be the one that **maximizes the likelihood** of observing these results.

Mathematical Formulation

Given data $X = \{x_1, x_2, \dots, x_n\}$ and a probability model $P(X|\theta)$ with parameters θ , we estimate θ by:

$$\theta^* = \arg \max_{\theta} P(X|\theta)$$

For independent data points:

$$L(\theta) = \prod_{i=1}^n P(x_i|\theta)$$

Taking the **log-likelihood** (to make computation easier):

$$\log L(\theta) = \sum_{i=1}^n \log P(x_i|\theta)$$

To find the best θ , we **differentiate** and solve:

$$\frac{d}{d\theta} \log L(\theta) = 0$$

Pseudo Code for MLE

plaintext

 Copy

 Edit

1. Define the probability distribution $P(X \mid \theta)$.
2. Compute the log-likelihood function.
3. Differentiate with respect to θ .
4. Solve for θ that maximizes the likelihood.

2. Maximum A Posteriori (MAP) Estimation

Intuition: Why Do We Need MAP?

MLE maximizes **only the likelihood**, ignoring any **prior knowledge**. MAP fixes this by incorporating **Bayes' Theorem**.

Example: Coin Flip Again

- Suppose we **already believe** the coin is biased towards heads.
 - MLE **ignores** this prior belief.
 - MAP **combines likelihood and prior**.
-

Mathematical Formulation

Using **Bayes' Rule**:

$$P(\theta|X) = \frac{P(X|\theta)P(\theta)}{P(X)}$$

MAP estimation **maximizes the posterior**:

$$\theta^* = \arg \max_{\theta} P(\theta|X)$$

Using **logarithms**:

$$\theta^* = \arg \max_{\theta} [\log P(X|\theta) + \log P(\theta)]$$

Difference from MLE?

- MLE: **Only uses** likelihood $P(X|\theta)$.
 - MAP: **Also considers** prior $P(\theta)$.
-

Pseudo Code for MAP

plaintext

Copy

Edit

1. Define the likelihood function $P(X | \theta)$.
2. Define the prior distribution $P(\theta)$.
3. Compute the posterior using Bayes' rule.
4. Maximize log posterior to find best θ .

3. Softmax Function

Intuition: Why Do We Need Softmax?

Softmax converts a **vector of arbitrary real numbers** into **probabilities** that sum to 1.

Example: Multi-Class Classification

Given class scores $[3.1, 2.5, 1.2]$, we want to **convert them into probabilities**.

Mathematical Formulation

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$

where:

- z_i is the **raw score** for class i .
 - $P(y_i)$ is the **probability** of class i .
-

Pseudo Code for Softmax

plaintext

Copy

Edit

1. Compute exponentials of all input scores.
2. Sum all exponentials.
3. Divide each exponentiated score by the sum.

4. Cross-Entropy Loss

Intuition: Why Do We Need Cross-Entropy?

Cross-entropy measures how different two probability distributions are. It is used as a loss function in classification.

Example: Classification

If the model predicts $P(\text{cat}) = 0.9$, and the correct label is **cat**, we want the loss to be **small**. If it predicts $P(\text{dog}) = 0.9$, we want the loss to be **high**.

Mathematical Formulation

For binary classification:

$$H(p, q) = - \sum p_i \log q_i$$

where:

- p is the **true probability** (ground truth).
- q is the **predicted probability**.

For multi-class:

$$L = - \sum_{i=1}^n y_i \log \hat{y}_i$$

Pseudo Code for Cross-Entropy

plaintext

 Copy

 Edit

1. Compute $-\log$ of predicted probability for the correct class.
2. Sum across all training examples.
3. Take the average.

1. Kullback-Leibler (KL) Divergence

Intuition: What Does KL Divergence Do?

KL Divergence measures how different one probability distribution is from another.

Example Intuition

- Suppose we have two distributions:
 - $P(x)$ (true distribution)
 - $Q(x)$ (our model's approximation)
 - KL divergence tells us **how much information is lost** when we approximate $P(x)$ using $Q(x)$.
-

Mathematical Formulation

For discrete distributions:

$$D_{\text{KL}}(P||Q) = \sum_i P(x_i) \log \frac{P(x_i)}{Q(x_i)}$$

For continuous distributions:

$$D_{\text{KL}}(P||Q) = \int P(x) \log \frac{P(x)}{Q(x)} dx$$

- If $P(x) = Q(x)$, KL divergence is **zero**.
- KL divergence is not symmetric: $D_{\text{KL}}(P||Q) \neq D_{\text{KL}}(Q||P)$.

Pseudo Code for KL Divergence

plaintext

 Copy

 Edit

1. Take two probability distributions P and Q .
2. Compute $\log(P / Q)$ for each value.
3. Multiply by P .
4. Sum over all values.

2. Jensen-Shannon Divergence (JSD)

Intuition: Why Use JSD Instead of KL Divergence?

- KL divergence is **not symmetric**.
 - JSD **fixes this by averaging KL divergences** between two distributions and a midpoint distribution.
-

Mathematical Formulation

Define the midpoint distribution:

$$M(x) = \frac{1}{2} (P(x) + Q(x))$$

JSD is:

$$D_{\text{JS}}(P||Q) = \frac{1}{2} D_{\text{KL}}(P||M) + \frac{1}{2} D_{\text{KL}}(Q||M)$$

- Unlike KL divergence, **JSD is symmetric**.
- JSD is **bounded between 0 and 1**, making it more stable.

Pseudo Code for Jensen-Shannon Divergence

plaintext

 Copy

 Edit

1. Compute $M = (P + Q) / 2$.
2. Compute $KL(P || M)$ and $KL(Q || M)$.
3. Take the average of both KL divergences.

3. Bayesian Inference

Intuition: Why Use Bayesian Inference?

- Unlike MLE, Bayesian inference **updates beliefs** as more data is observed.
 - Instead of a single **best** parameter value, it provides a **probability distribution over parameters**.
-

Mathematical Formulation (Bayes' Theorem)

$$P(\theta|X) = \frac{P(X|\theta)P(\theta)}{P(X)}$$

where:

- $P(\theta|X)$ → **Posterior** (updated belief about θ after seeing data)
 - $P(X|\theta)$ → **Likelihood** (how well θ explains data)
 - $P(\theta)$ → **Prior** (belief about θ before seeing data)
 - $P(X)$ → **Evidence** (normalizing constant)
-

Pseudo Code for Bayesian Inference

plaintext

Copy

Edit

1. Define prior $P(\theta)$.
2. Compute likelihood $P(X | \theta)$.
3. Compute posterior $P(\theta | X)$ using Bayes' rule.



4. Gaussian Mixture Models (GMMs)

Intuition: Why Use GMM?

- K-Means assumes clusters are spherical (equal variance).
 - GMM allows clusters of different shapes by using multiple Gaussian distributions.
-

Mathematical Formulation

A GMM models data as a weighted sum of multiple Gaussians:

$$P(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x | \mu_k, \Sigma_k)$$

where:

- $\pi_k \rightarrow$ Mixing coefficient (probability of cluster k)
 - $\mathcal{N}(x | \mu_k, \Sigma_k) \rightarrow$ Gaussian distribution for cluster k
 - $K \rightarrow$ Number of clusters
-

Pseudo Code for GMM

plaintext

Copy

Edit

```
1. Initialize means, covariances, and mixing coefficients.
2. E-Step: Compute probabilities of each point belonging to each cluster.
3. M-Step: Update parameters using weighted averages.
4. Repeat until convergence.
```



5. Expectation-Maximization (EM Algorithm)

Intuition: Why Use EM?

- When data has **missing values** or **latent variables**, direct optimization is difficult.
 - EM iteratively estimates hidden variables and optimizes parameters.
-

Mathematical Formulation

1. **E-Step:** Compute the **expected value** of hidden variables.
2. **M-Step:** Maximize the expected likelihood.

For parameter θ :

$$Q(\theta|\theta^{(t)}) = \mathbb{E} \left[\log P(X, Z|\theta) | X, \theta^{(t)} \right]$$

Update:

$$\theta^{(t+1)} = \arg \max_{\theta} Q(\theta|\theta^{(t)})$$

Expectation-Maximization (EM) Algorithm – Pseudo Code with Explicit Math

1. Initialize parameters:

$$\theta^{(0)} = \{\pi_k, \mu_k, \Sigma_k\} \quad (\text{randomly initialized})$$

2.

Repeat until convergence:

Step 1: Expectation Step (E-Step)

Compute the **responsibilities** (posterior probability of cluster k for each data point i):

$$\gamma_{ik} = \frac{\pi_k^{(t)} \mathcal{N}(x_i | \mu_k^{(t)}, \Sigma_k^{(t)})}{\sum_{j=1}^K \pi_j^{(t)} \mathcal{N}(x_i | \mu_j^{(t)}, \Sigma_j^{(t)})}$$

where $\mathcal{N}(x | \mu, \Sigma)$ is the Gaussian probability density function:

$$\mathcal{N}(x | \mu, \Sigma) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$

Step 2: Maximization Step (M-Step)

Update the **mixing coefficients**:

$$\pi_k^{(t+1)} = \frac{1}{N} \sum_{i=1}^N \gamma_{ik}$$

Update the **means**:

$$\mu_k^{(t+1)} = \frac{\sum_{i=1}^N \gamma_{ik} x_i}{\sum_{i=1}^N \gamma_{ik}}$$

Update the **covariances**:

$$\Sigma_k^{(t+1)} = \frac{\sum_{i=1}^N \gamma_{ik} (x_i - \mu_k^{(t+1)})(x_i - \mu_k^{(t+1)})^T}{\sum_{i=1}^N \gamma_{ik}}$$

Step 3: Compute Log-Likelihood

Compute the **log-likelihood** to track convergence:

$$L(\theta^{(t)}) = \sum_{i=1}^N \log \sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(x_i | \mu_k^{(t)}, \Sigma_k^{(t)})$$

Step 4: Check Convergence

If the **change in log-likelihood** is below a threshold ϵ :

$$|L(\theta^{(t+1)}) - L(\theta^{(t)})| < \epsilon$$

then **stop the iterations**.

Key Takeaways

- ✓ The E-step computes expected assignments of data points to clusters (soft clustering).
- ✓ The M-step updates the model parameters using these expectations.
- ✓ Convergence is determined by log-likelihood improvement.
- ✓ Used in Gaussian Mixture Models (GMMs), Hidden Markov Models (HMMs), and other latent variable models.

